



CHALMERS
UNIVERSITY OF TECHNOLOGY



UNIVERSITY OF GOTHENBURG

DAT490 Introduction

Dr. Philipp Leitner



philipp.leitner@chalmers.se



@xLeitix

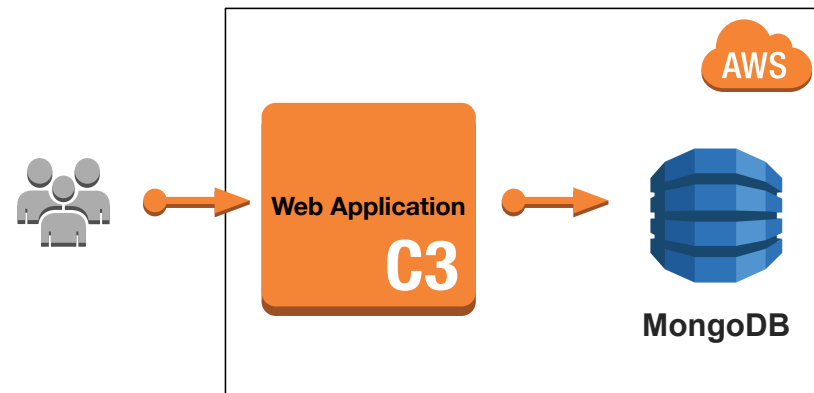
LECTURE 1

Covers ...

**A gentle introduction to Architecting for Scale
Scalability and Availability**

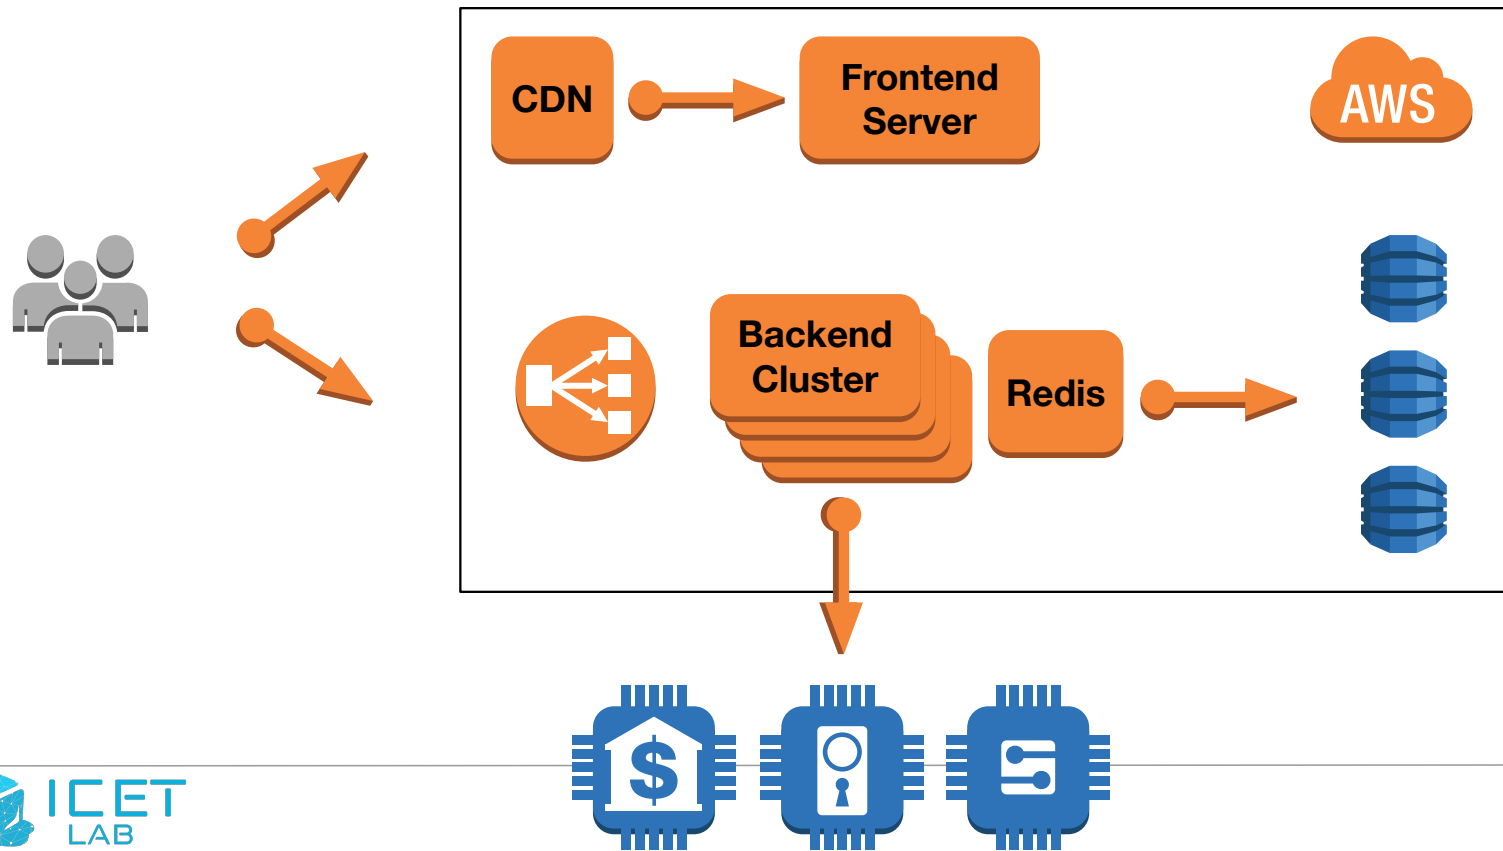
Architectural Evolution of a Web App

Stage 1 - a three-tier web application



Architectural Evolution of a Web App

Stage 2 - architecting for (reasonable) scale





CHALMERS
UNIVERSITY OF TECHNOLOGY




UNIVERSITY OF GOTHENBURG

■ ■ ■ ■

Architectural Evolution of a Web App

Stage n - microservices



Source: amazon.com

What drives this evolution?

What drives this evolution?

Support for new features (*of course*)

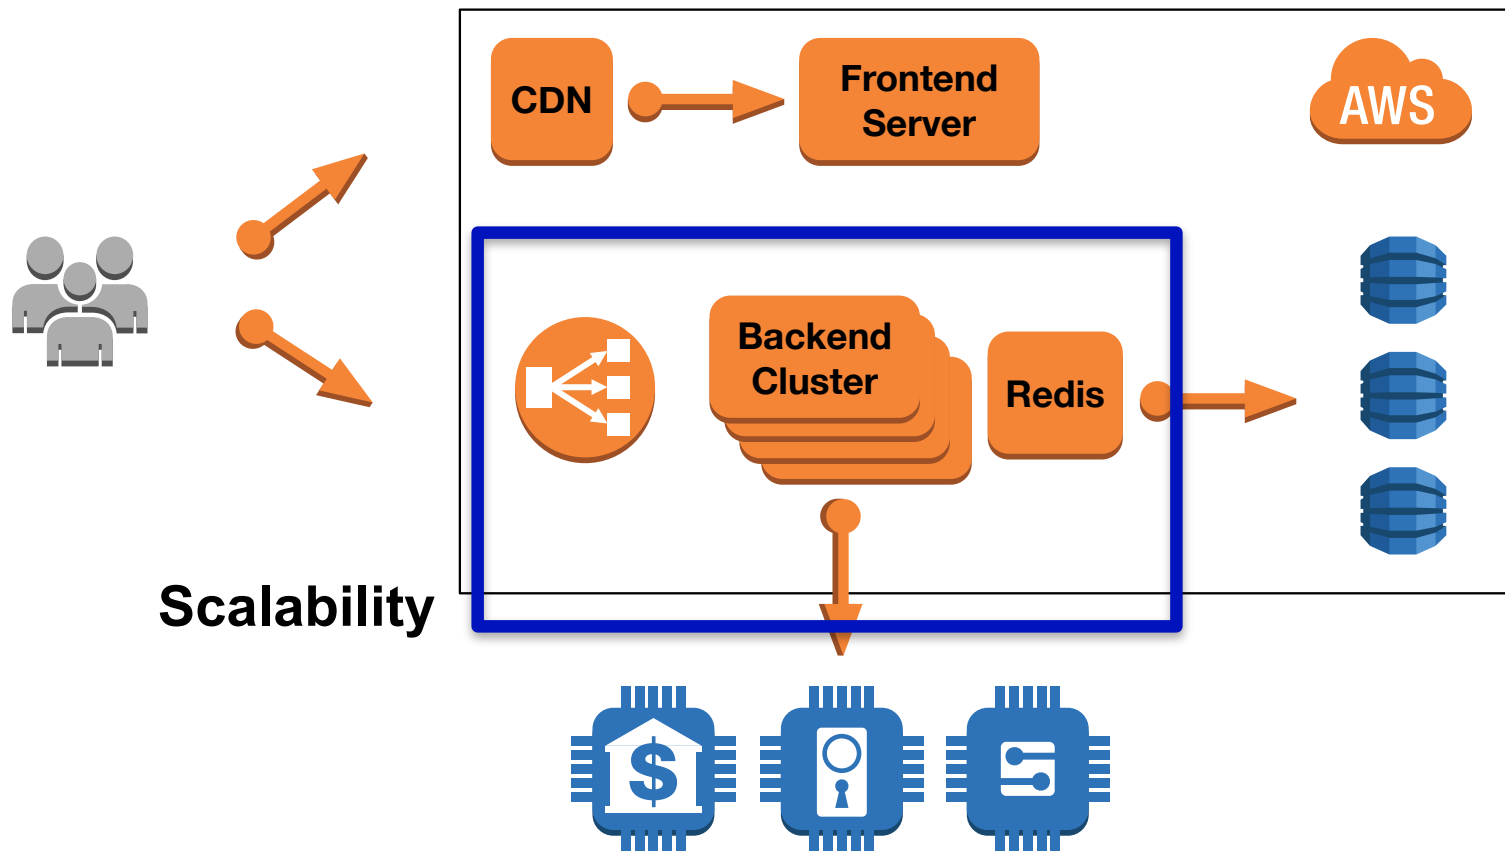
What drives this evolution?

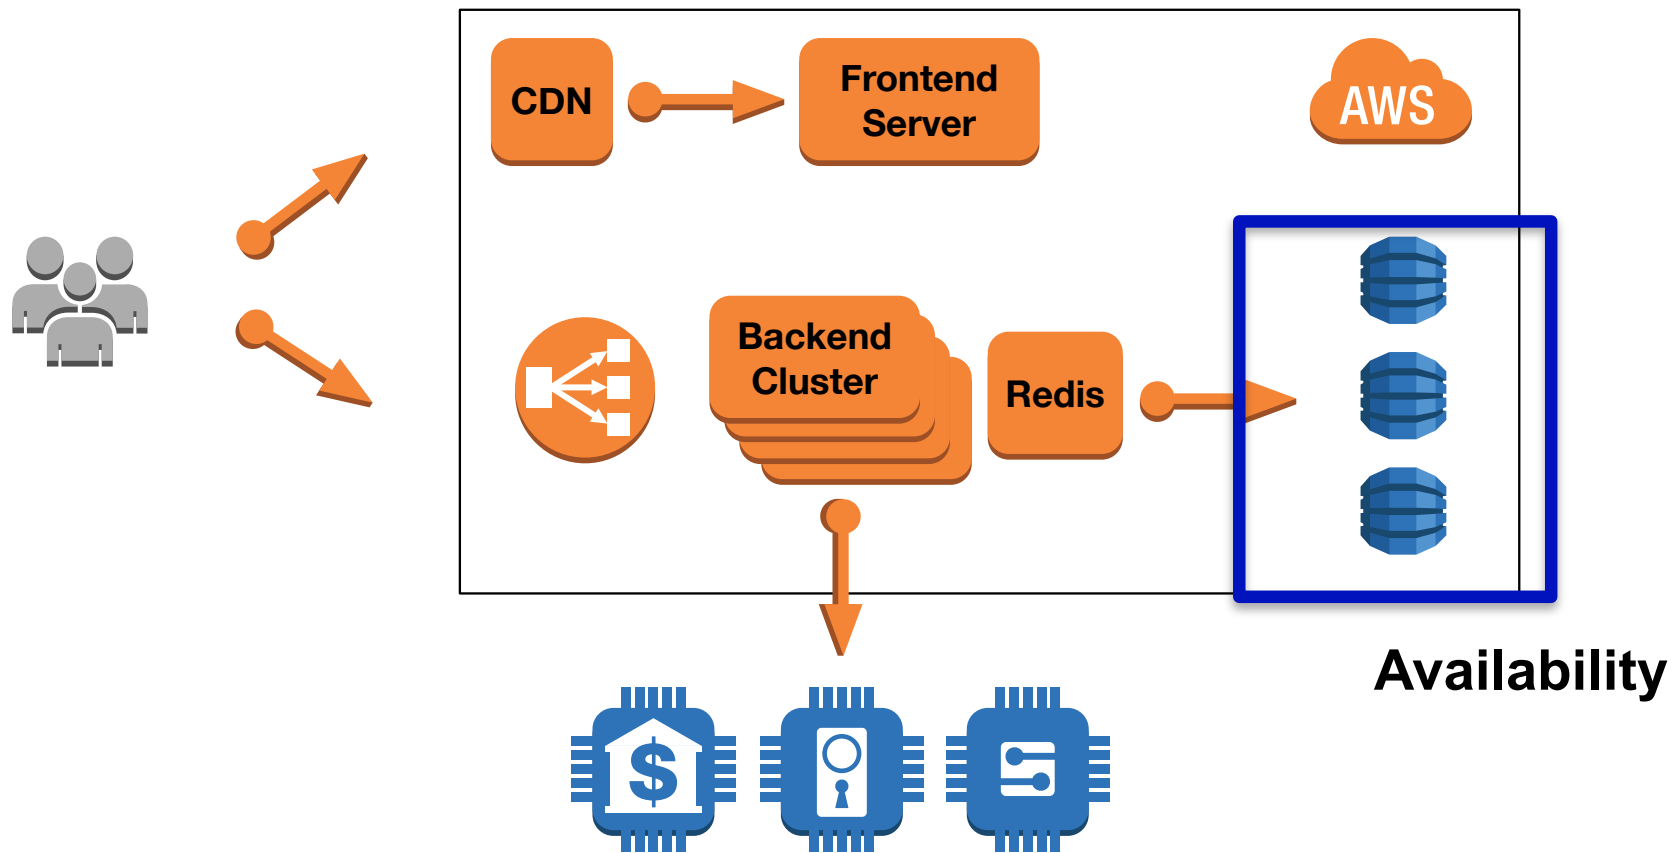
Support for new features (*of course*)


But also:

Scalability (performance at scale, scaling development)

Availability (resilience, survivability)







Architecting for Scale

*“Transforming a (simple) application into a **system** that can be used by a large number of concurrent users (**scalability**), that can be maintained and extended by a large team (**scalability of development**), that requires low or no downtimes (**availability**), and which survives intermittent failures (**resilience**).”*

Scalability (in terms of “number of concurrent users”)

Scalability is the property of a system to handle a growing amount of work by adding resources to the system.

Metric for scalability:

Speedup

**how much does my system become faster
when adding 1 new unit of computing
resources?**

$$\textit{Speedup} = \frac{N_{orig} t_{new}}{t_{orig} N_{new}}$$

e.g.

$$\textit{Speedup} = \frac{1 * 80}{100 * 2} = 0.4$$

Speedup

Speedup is in $[0;1]$

Special cases:

Speedup = 0

no gain, e.g., adding more CPUs to a single-threaded program

Speedup = 1

Twice the resources, half the time (perfect speedup)

“Embarrassingly parallel” applications

Amdahl's Law


$$t_{total} = t_{serial} + \frac{t_{parallel}}{N}$$

$$Speedup = \frac{1}{(1 + (N - 1)a)}, \text{ with } a = \frac{t_{serial}}{(t_{serial} + t_{parallel})}$$

For $a > 0$, speedup converges against 0 (quickly)



Amdahl's Law



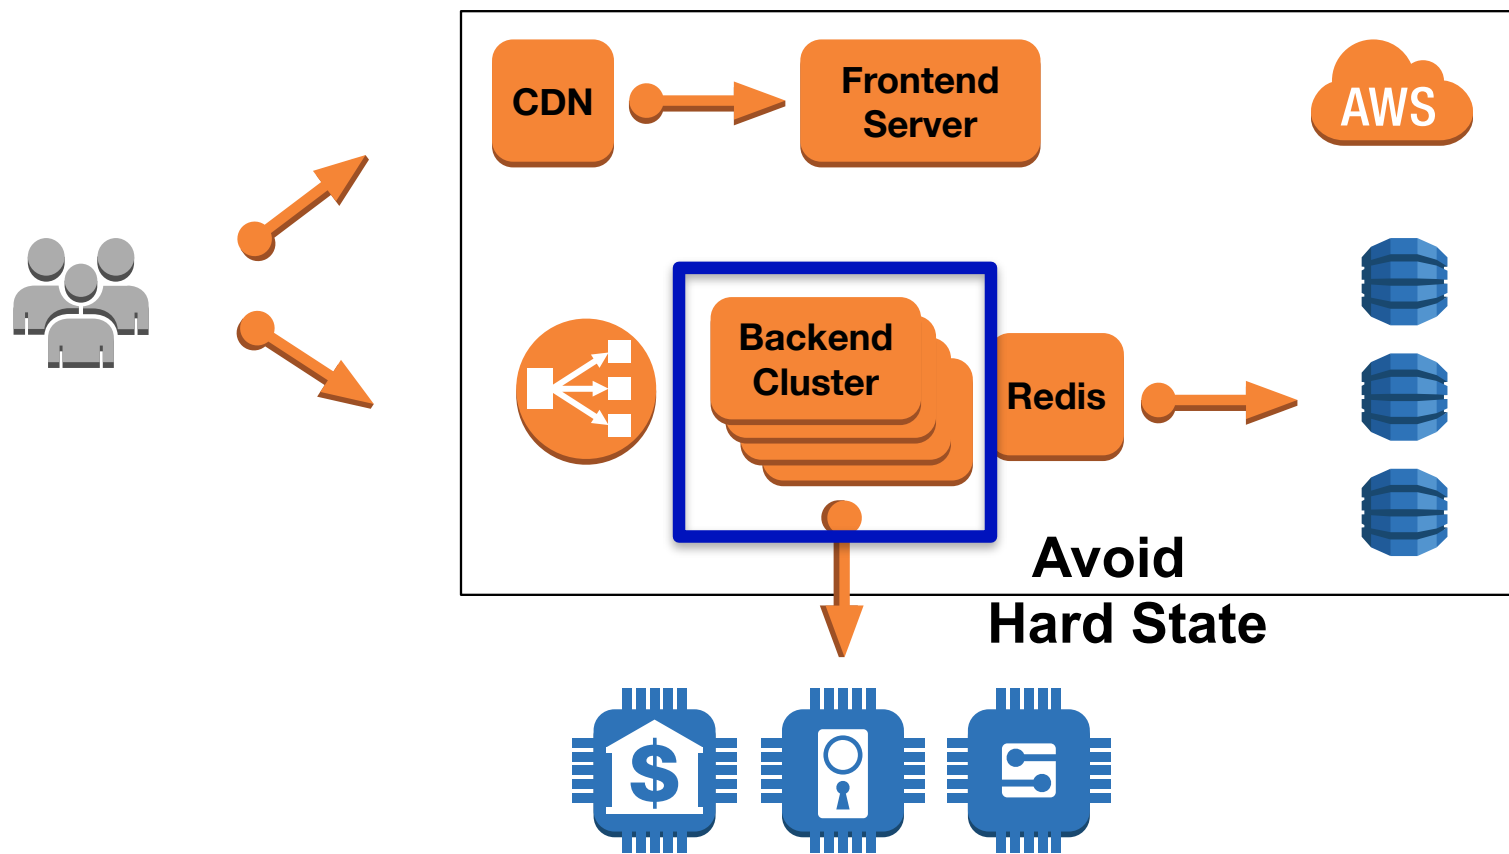
Building Embarrassingly Parallel Web Apps

Implication:

State (another word for “serial processing required”) is the archenemy of architecting for scale

Most architectural guidelines for building scalable Web applications have some “statelessness” requirement:

e.g., the REST architectural style, Heroku’s 12-Factor App, etc.



Performance

Page Load Time

Time from request sent to page completely loaded

Response Time

Time from request sent to response received

Network Latency

Time request / response takes for transmission through the network

Throughput

Number of requests that can be processed per time unit

Error Rate

How frequently to requests transiently fail?

For all these:

Average / 95-percentile / 99-percentile

#RESPONSETIME



Performance vs. Scalability in Practice

In Web development practice, the relationship between performance and scalability is ... complicated.

... looking at Amdahl's law, we would expect that twice as many backend instances would lead to half the response time

Spoiler: this is never how it actually works

Availability

Availability is the probability of a system to be online (“available”) at any point in time

$$av(S) = 1 - \frac{downtime}{uptime}$$

Availability

Common measure for availability: x-nines,

“2-nines” (0.99) ... about 8 hours offline per month

“3-nines” (0.999) ... about 44 minutes offline per month

“5-nines” (0.99999) ... about half a minute offline per month

Service Level Agreements (SLAs)

Service providers often contractually guarantee a certain availability (and sometimes response time)

E.g., Google:

“

Bigtable instances with a multi-cluster routing policy across three or more regions are now covered by a 99.999% monthly uptime percentage under the new SLA. Bigtable supports 99.99% monthly uptime percentage for all instances with a multi-cluster routing policy across less than three regions and 99.9% monthly uptime percentage for all instances with a Single-Cluster routing policy.

Service Level Agreements (SLAs)

Important questions for SLAs:

How is the metric specified (exactly), and who measures it?

Especially, what's the aggregation time window?

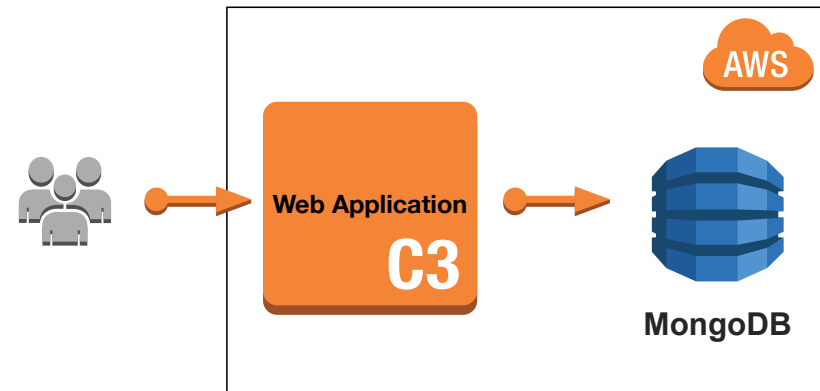
What's the compensation for SLA violations?

Do scheduled maintenance windows count?

Etc.

Redundancy

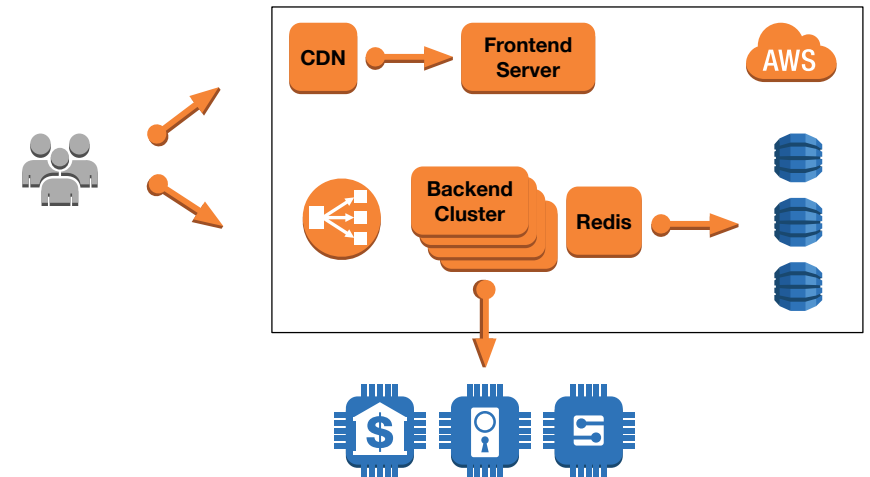
One (common) way to improve availability is **redundancy**
(Avoiding single points of failure)



Availability of our simple architecture is:

$$av = av(app) * av(db)$$

$$av = 0.99 * 0.99 = 0.98$$



Availability of our “scaled” architecture:

$av = av(cdn) * av(any_backend) * av(any_db)$

$av \sim 1$ (theoretically)

Co-Located Failures

In practice our availability will still be < 1
Mostly due to “co-located failures” (co-failures)

Availabilities of our services are in practice not independent
One service failing **drastically** increases the failure probability of other services

E.g., if all our services are in one AWS data center and this data center burns down, it matters little *how many* replicas we had

In practice availability of > 0.99999 is not consistently achievable

For the third time in less than a month **Amazon Web Services (AWS)** is down, not only affecting its own store, but third party services that host with them like Epic Games Store, Tinder, and Hulu among others.

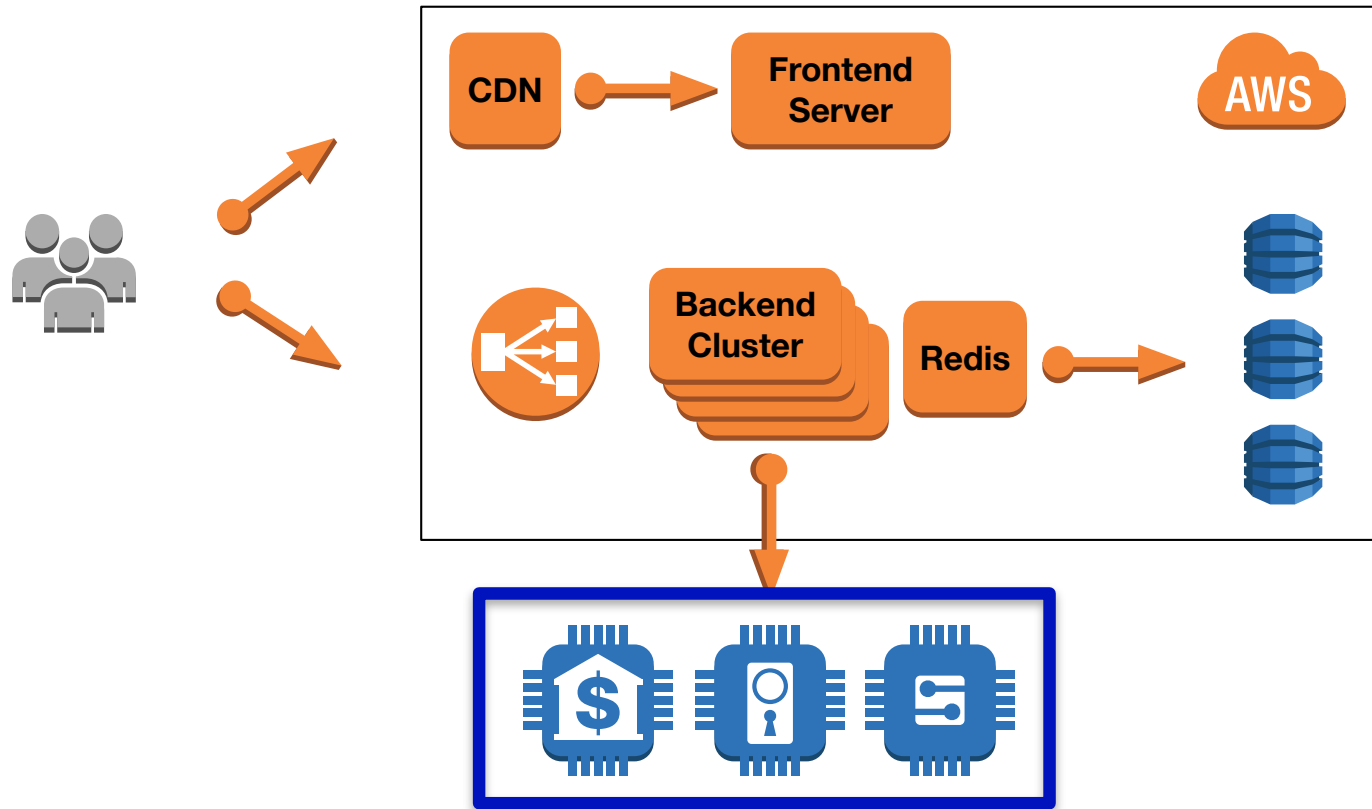
<https://www.marca.com/en/lifestyle/us-news/2021/12/22/61c36eb5ca4741ee0f8b45ec.html>


Resilience

Another way to improve our availability is resilience (or **survivability)**

Essentially:

Make your application react gracefully to optional components not being available





Summary

Scalability and availability as driving the migration to scaled architectures

