



UNIVERSITY OF GOTHENBURG

Kubernetes

Dr. Philipp Leitner



philipp.leitner@chalmers.se



[@xLeitix](https://twitter.com/xLeitix)

LECTURE 4

Covers ...

Kubernetes

Helm (the Kubernetes Package Manager)

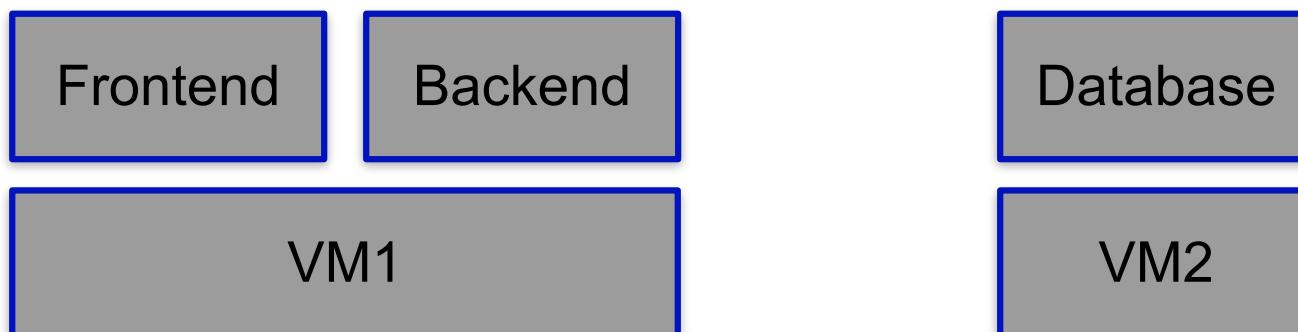
Note: much of this lecture is based on examples and slides by Lucas Käldström (@kubernetesonarm)

Kubernetes

<https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/>

Docker-compose allows you to easily start a composition **on a single computer**
... but that's not how we actually want to deploy our application in production

"Kubernetes is a portable, extensible, open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation."



What's Kubernetes?

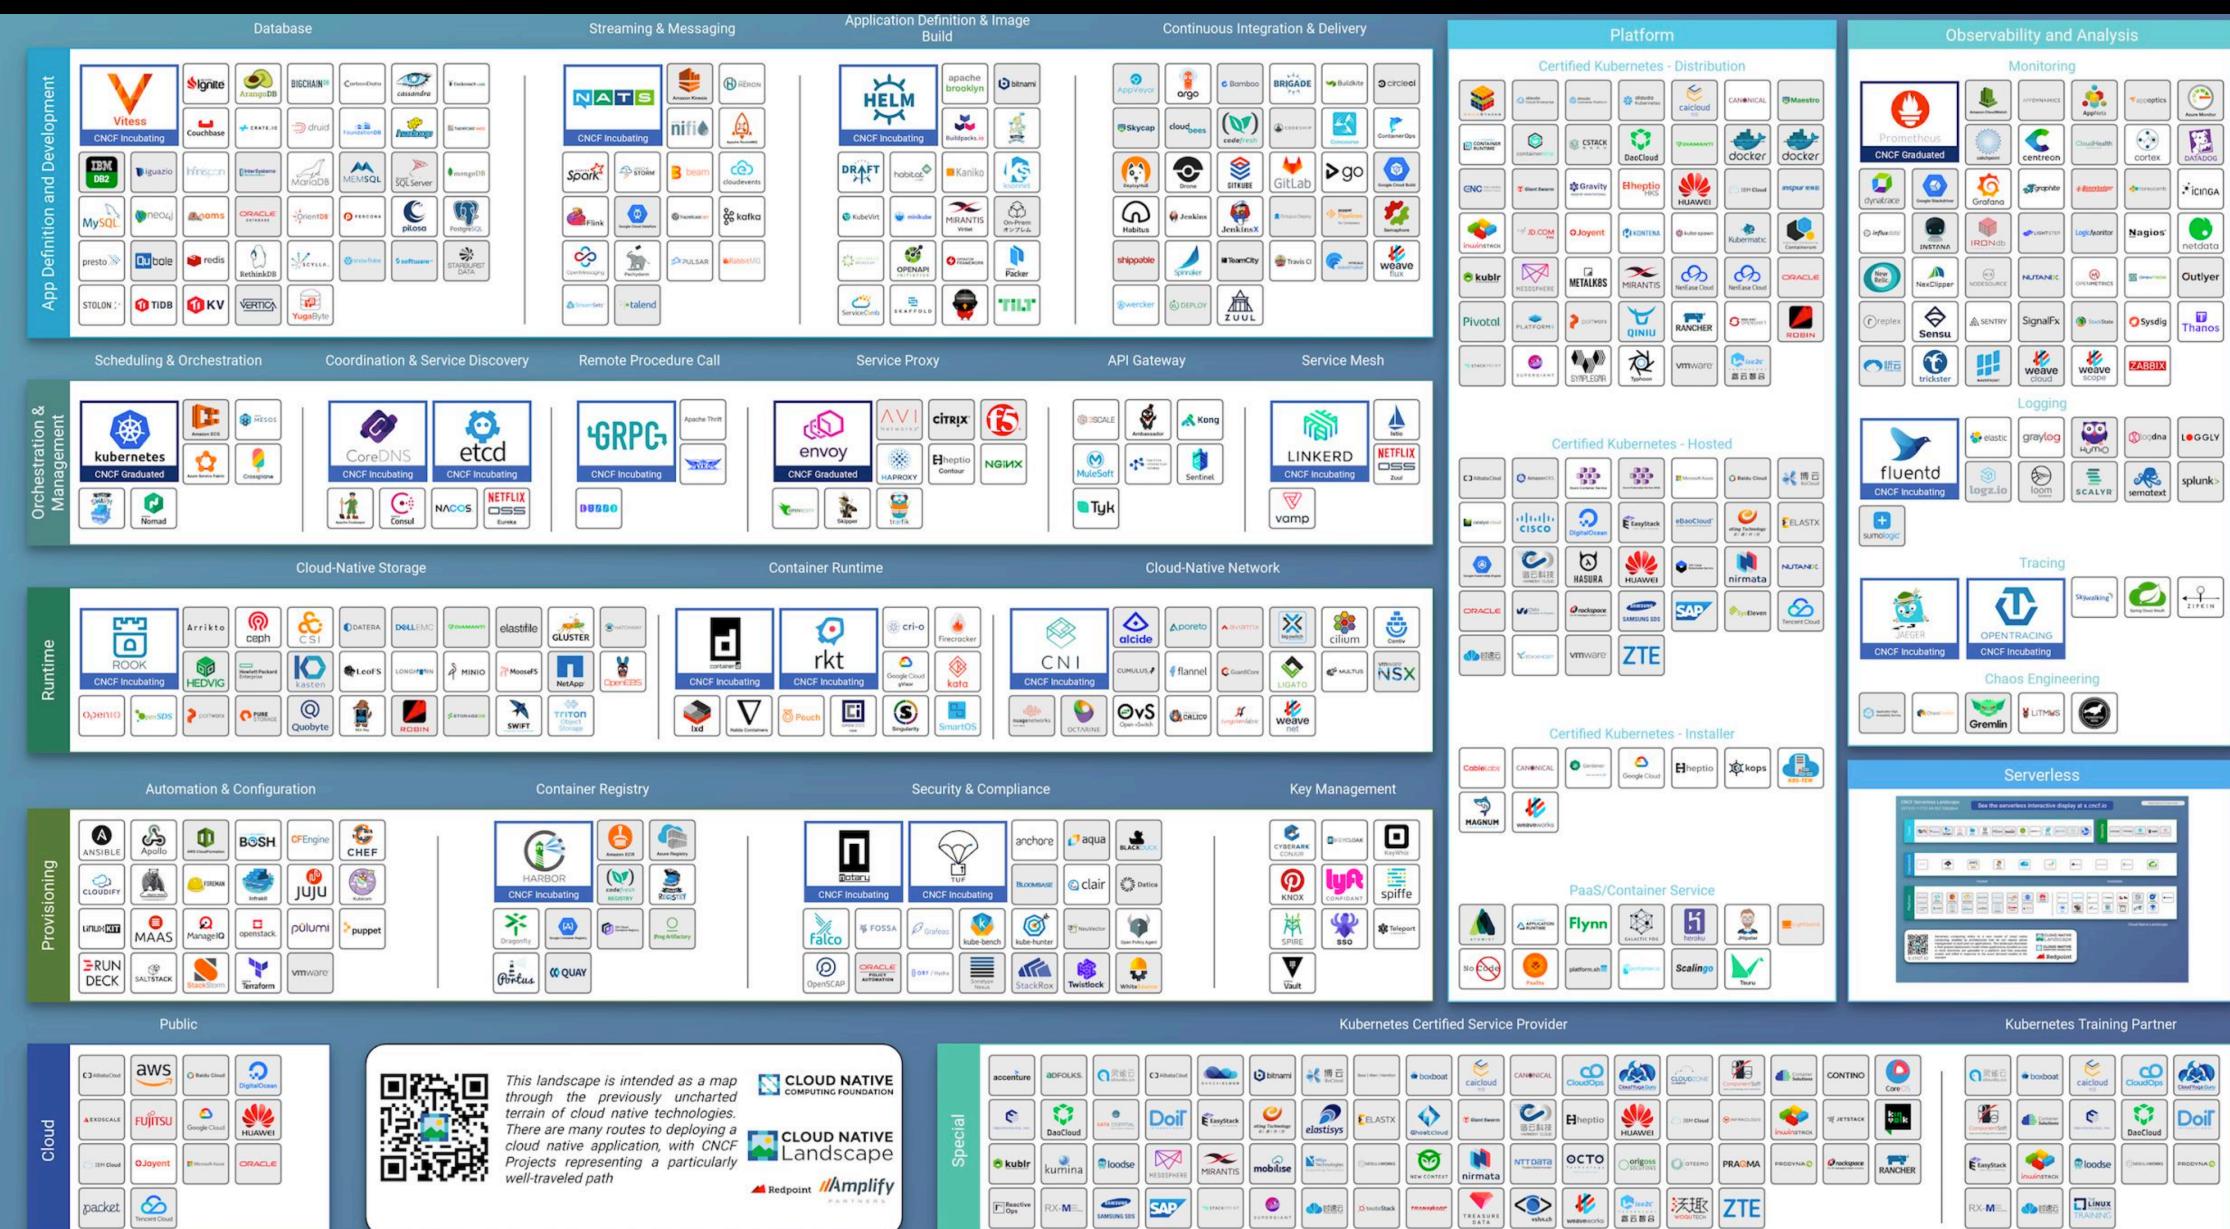
A Production-Grade Container Orchestration System

Open source project spun out of Google (“Project Borg”)

Developed within CNCF, which is itself part of the Linux foundation

Incredibly fancy right now (~ 2 billion containers launched per week)

Extremely large and complex ecosystem





What's Kubernetes?

Basic tenet of Kubernetes is that it's declarative and self-healing

Users declare the state of their application in config files

Kubernetes attempts to always keep the application in that state

What's Kubernetes?

Example 1: (availability)

User: "My application consists of 3 MongoDB shards"

k8s: starts three pods with MongoDB running

k8s: observes that one MongoDB shard crashed, starts another

Example 2: (scalability, auto-scaling)

User: "My application runs best with a CPU load [0.4;0.6]"

k8s: starts default number of pods, observes CPU utilization

k8s: observes that CPU utilization >0.6, starts more pods

k8s: observes that CPU utilization <0.4, releases some pods

Core Terminology

<https://kubernetes.io/docs/concepts/>

Node

A (physical or virtual) resources that Kubernetes should deploy the application *on*.
If you use a hosted Kubernetes cluster (like GKE) the resources (nodes) will be VMs provided by the cloud

Cluster

A set of nodes, managed by a Kubernetes installation

Control Plane

The management / scheduling functionality of the cluster, provided by the cluster provider (GCE in our case)

Pod

A group of one or multiple containers that should be deployed together, on the same node

Core Terminology

<https://kubernetes.io/docs/concepts/>

Service

A group of one or multiple pods that together forms something that can be accessed through the Internet, e.g., a Web application or API

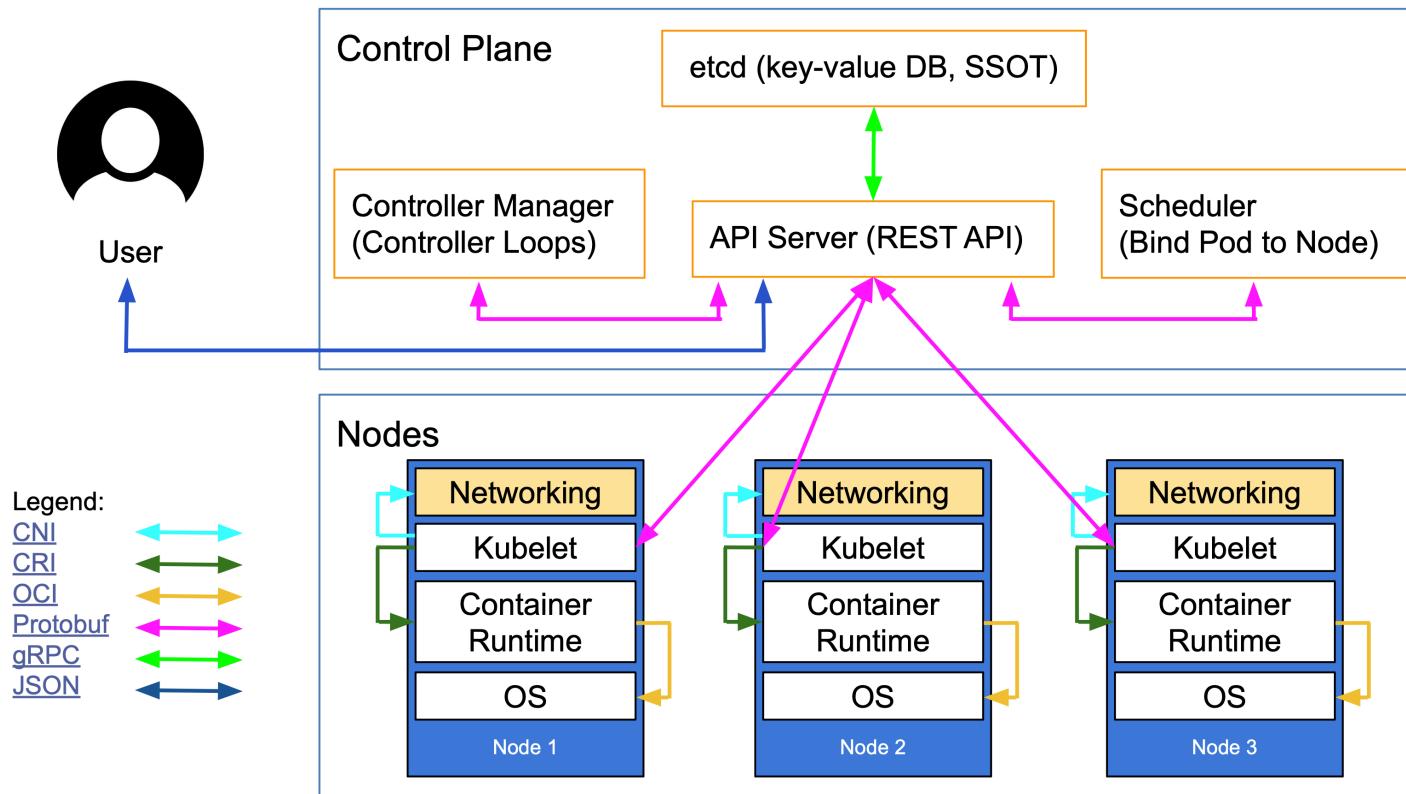
Workload

An application

Namespace

A way to group and isolate different workloads etc. within a cluster

High-Level Architecture



Getting Started with Kubernetes on GCE

Step 1: Create a cluster

```
gcloud container clusters create \
--region europe-north1-a --num-nodes 2 dat490
```

Getting Started with Kubernetes on GCE

Step 2:

Run an example job on the cluster

```
kubectl create job hello \
--image=busybox -- echo "Hello World"
```

Getting Started with Kubernetes on GCE

Step 3: using our cluster

Pods, deployments, services, ingress

Core components of a Kubernetes application

Defined through YAML files (manifest files)

Creating a Pod

Atomic deployable unit

(c) Lucas Käldström

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
  labels:
    app: nginx
spec:
  containers:
  - image: nginx:1.13.9
    name: nginx
  ports:
  - name: http
    containerPort: 80
```

Creating a Pod

Register manifest with cluster:

```
kubectl apply -f pod.yaml
```

Creating a Deployment

Essentially a collection of pods

(c) Lucas Käldström

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: nginx
    name: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - image: nginx:1.13.9
        name: nginx
      ports:
      - name: http
        containerPort: 80
```

Pod Template

Creating a Deployment

Register manifest with cluster:

```
kubectl apply -f deployment.yaml
```

Creating a Service

Provides a “front-end” for a set of pods
(often a deployment)

Accessible within the cluster through
Internal IP address or internal DNS name:

nginx.default.svc.cluster.local

(c) Lucas Käldström

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
  namespace: default
  labels:
    app: nginx
spec:
  type: ClusterIP
  ports:
  - name: http
    port: 80
    targetPort: 80
  selector:
    app: nginx
```

Pod Selector

Creating a Deployment

Register manifest with cluster:

```
kubectl apply -f service.yaml
```

Send a test request:

```
kubectl create job testnginx\  
--image=busybox -- wget http://nginx.default.svc.cluster.local
```

(check logs of created pod)

Ingress

Ingress is a separate service that acts as the gateway between the “outer world” and a service.

Separate tool that should either exist externally or be installed in the cluster.

Path of an external request:

Request -> Ingress -> Service -> one pod in the deployment

(c) Lucas Käldström

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: nginx
  namespace: default
  labels:
    app: nginx
spec:
  rules:
  - host: nginx.kubernetesfinland.com
    http:
      paths:
      - path: /
        backend:
          serviceName: nginx
          servicePort: 80
```

Service Reference

Deployment without Explicit Ingress

Configure the type “LoadBalancer” rather than “ClusterIP” in `service.yaml`

```
service:  
  type: LoadBalancer  
  port: 80  
  targetPort: 80
```



Helm

So far we have seen how to configure k8s “manually”
(through a large number of increasingly complex YAML files)

Helm (<https://helm.sh>) is the standard **package manager** for Kubernetes

Two core use cases:

- (1) Easily install existing workloads (applications) in our cluster
- (2) Package our own applications using Helm and deploy all in one go (rather than registering manifest by manifest)

Using Helm to Install an Existing Application

Let's install MongoDB as an example workload via Helm
(also part of Assignment 2)

Add package repository:

```
helm repo add bitnami\  
https://charts.bitnami.com/bitnami
```

Install package:

```
helm install mongo bitnami/mongodb
```

values.yaml

Typically, Helm packages (“Helm charts”) can be configured / customised through a **values.yaml file**

```
helm install -f values.yaml\  
gitlab-runner gitlab/gitlab-runner
```

values.yaml

Configurations can also be applied retroactively

```
helm upgrade -f values.yaml\  
gitlab-runner gitlab/gitlab-runner
```

Packaging our own application as a Helm Chart

Helm charts are written as **templates** using the Go templating system **tpl**

Templates for the different k8s manifest files that are needed to deploy the app

Creating a new chart:

```
helm create my-application
```

```
phileitn@dev-machine:~$ ls my-application/
Chart.yaml  charts  templates  values.yaml
```

Structure of a Helm Package

Chart.yaml ... basic information about the package

values.yaml ... default values, can be overwritten during install

templates (folder) ... templates for all types of manifests that a k8s application tends to need (deployment, service, ingress, ...)

templates/NOTES.txt ... what gets outputted after the install

Packaging and installing

Lint (check) our chart

```
helm lint my-application
```

Create the package

```
helm package my-application
```

Install on cluster

```
helm install my-application my-application-0.1.0.tgz
```